

## Notebook:

- [Link](#)

## Cluster:

- *N2-highmem-4, 2 - 8 autoscaling. 16.4 LTS*

## Query Without Materialization:

### Query:

```
WITH joined_files AS (  
  SELECT  
    f.repo_name,  
    f.ref,  
    f.path,  
    f.mode,  
    f.id,  
    f.symlink_target,  
    c.size,  
    c.content,  
    c.binary,  
    c.copies  
  FROM  
    files f  
    INNER JOIN contents c  
      ON f.id = c.id  
)  
SELECT  
  repo_name,  
  path,  
  size,  
  RANK() OVER (PARTITION BY repo_name ORDER BY size DESC) AS size_rank,  
  copies,  
  binary
```

FROM

joined\_files

ORDER BY

repo\_name,

size\_rank;

Explain Plan:

```
== Physical Plan ==
AdaptiveSparkPlan (26)
+- == Initial Plan ==
   ColumnarToRow (25)
   +- PhotonResultStage (24)
      +- PhotonSort (23)
         +- PhotonShuffleExchangeSource (22)
            +- PhotonShuffleMapStage (21)
               +- PhotonShuffleExchangeSink (20)
                  +- PhotonWindow (19)
                     +- PhotonSort (18)
                        +- PhotonShuffleExchangeSource (17)
                           +- PhotonShuffleMapStage (16)
                              +- PhotonShuffleExchangeSink (15)
                                 +- PhotonProject (14)
                                    +- PhotonShuffledHashJoin Inner (13)
                                       :- PhotonShuffleExchangeSource (6)
                                       : +- PhotonShuffleMapStage (5)
                                       : +- PhotonShuffleExchangeSink
(4)
                                       :
                                       : +- PhotonFilter (3)
                                       : +- PhotonRowToColumnar
(2)
                                       :
                                       +- Scan
BigQueryV1RelationFromV2ScanWithNoMaterialization(org.apache.spark.sql.SQLContext@934053c,StructType(StructField(repo_name,StringType,true),StructField(path,StringType,true),StructField(id,StringType,true)),BigQueryScanPushDowns(Some(StructType(StructField(repo_name,StringType,true),StructField(path,StringType,true),StructField(id,StringType,true))),[Long.apache.spark.sql.connector.expressions.filter.Predicate;@5e71c3dd,None,None,[Ljava.lang.String;@7b8a780a,[Ljava.lang.String;@7b8a780a,[Ljava.lang.String;@7b8a780a,[Long.apache.spark.sql.sources.Filter;@208d329b),DirectBigQueryRelation[fe-dev-
```

```

sandbox.dkushari_ds.files]) dkushari_bq_test.dkushari_ds.files (1)
    +- PhotonShuffleExchangeSource (12)
      +- PhotonShuffleMapStage (11)
        +- PhotonShuffleExchangeSink
(10)
    +- PhotonFilter (9)
      +- PhotonRowToColumnar
(8)
    +- Scan
BigQueryV1RelationFromV2ScanWithNoMaterialization(org.apache.spark.sql.SQLC
ontext@934053c,StructType(StructField(id,StringType,true),StructField(size,
LongType,true),StructField(binary,BooleanType,true),StructField(copies,Long
Type,true)),BigQueryScanPushDowns(Some(StructType(StructField(id,StringType
,true),StructField(size,LongType,true),StructField(binary,BooleanType,true)
,StructField(copies,LongType,true))),[Long.apache.spark.sql.connector.expre
ssions.filter.Predicate;@1b872ece,None,None,[Ljava.lang.String;@7b8a780a,[L
java.lang.String;@7b8a780a,[Ljava.lang.String;@7b8a780a,[Long.apache.spark.
sql.sources.Filter;@73415c2d),DirectBigQueryRelation[fe-dev-sandbox.dkushar
i_ds.contents]) dkushari_bq_test.dkushari_ds.contents (7)

```

(1) Scan

```

BigQueryV1RelationFromV2ScanWithNoMaterialization(org.apache.spark.sql.SQLC
ontext@934053c,StructType(StructField(repo_name,StringType,true),StructFiel
d(path,StringType,true),StructField(id,StringType,true)),BigQueryScanPushDo
wns(Some(StructType(StructField(repo_name,StringType,true),StructField(path
,StringType,true),StructField(id,StringType,true))),[Long.apache.spark.sql.
connector.expressions.filter.Predicate;@5e71c3dd,None,None,[Ljava.lang.Stri
ng;@7b8a780a,[Ljava.lang.String;@7b8a780a,[Ljava.lang.String;@7b8a780a,[Lor
g.apache.spark.sql.sources.Filter;@208d329b),DirectBigQueryRelation[fe-dev-
sandbox.dkushari_ds.files]) dkushari_bq_test.dkushari_ds.files

```

Output [3]: [repo\_name#217, path#219, id#221]

Arguments: [repo\_name#217, path#219, id#221],

```

[StructField(repo_name,StringType,true), StructField(path,StringType,true),
StructField(id,StringType,true)],

```

PushedDownOperators(None,None,None,None,List(),List()),

PreScala213BigQueryRDD[54] at RDD at PreScala213BigQueryRDD.java:40,

```

BigQueryV1RelationFromV2ScanWithNoMaterialization(org.apache.spark.sql.SQLC
ontext@934053c,StructType(StructField(repo_name,StringType,true),StructFiel
d(path,StringType,true),StructField(id,StringType,true)),BigQueryScanPushDo
wns(Some(StructType(StructField(repo_name,StringType,true),StructField(path
,StringType,true),StructField(id,StringType,true))),[Long.apache.spark.sql.
connector.expressions.filter.Predicate;@5e71c3dd,None,None,[Ljava.lang.Stri

```

```
ng;@7b8a780a,[Ljava.lang.String;@7b8a780a,[Ljava.lang.String;@7b8a780a,[Long.apache.spark.sql.sources.Filter;@208d329b),DirectBigQueryRelation[fe-dev-sandbox.dkushari_ds.files]), `dkushari_bq_test`.`dkushari_ds`.`files`, Statistics(sizeInBytes=8.0 EiB, ColumnStat: N/A)
```

(2) PhotonRowToColumnar

Input [3]: [repo\_name#217, path#219, id#221]

(3) PhotonFilter

Input [3]: [repo\_name#217, path#219, id#221]

Arguments: isnotnull(id#221)

(4) PhotonShuffleExchangeSink

Input [3]: [repo\_name#217, path#219, id#221]

Arguments: hashpartitioning(id#221, 200)

(5) PhotonShuffleMapStage

Input [3]: [repo\_name#217, path#219, id#221]

Arguments: ENSURE\_REQUIREMENTS, [id=#629]

(6) PhotonShuffleExchangeSource

Input [3]: [repo\_name#217, path#219, id#221]

(7) Scan

```
BigQueryV1RelationFromV2ScanWithNoMaterialization(org.apache.spark.sql.SQLContext@934053c,StructType(StructField(id,StringType,true),StructField(size,LongType,true),StructField(binary,BooleanType,true),StructField(copies,LongType,true)),BigQueryScanPushDowns(Some(StructType(StructField(id,StringType,true),StructField(size,LongType,true),StructField(binary,BooleanType,true),StructField(copies,LongType,true))),[Long.apache.spark.sql.connector.expressions.filter.Predicate;@1b872ece,None,None,[Ljava.lang.String;@7b8a780a,[Ljava.lang.String;@7b8a780a,[Ljava.lang.String;@7b8a780a,[Long.apache.spark.sql.sources.Filter;@73415c2d),DirectBigQueryRelation[fe-dev-sandbox.dkushari_ds.contents]) dkushari_bq_test.dkushari_ds.contents
```

Output [4]: [id#228, size#229L, binary#231, copies#232L]

Arguments: [id#228, size#229L, binary#231, copies#232L],

[StructField(id,StringType,true), StructField(size,LongType,true), StructField(binary,BooleanType,true), StructField(copies,LongType,true)],

PushedDownOperators(None,None,None,None,List(),List()),

PreScala213BigQueryRDD[55] at RDD at PreScala213BigQueryRDD.java:40,

```
BigQueryV1RelationFromV2ScanWithNoMaterialization(org.apache.spark.sql.SQLContext@934053c,StructType(StructField(id,StringType,true),StructField(size,LongType,true),StructField(binary,BooleanType,true),StructField(copies,Long
```

```
Type,true)),BigQueryScanPushDowns(Some(StructType(StructField(id,StringType,true),StructField(size,LongType,true),StructField(binary,BooleanType,true),StructField(copies,LongType,true))),[Long.apache.spark.sql.connector.expressions.filter.Predicate;@1b872ece,None,None,[Ljava.lang.String;@7b8a780a,[Ljava.lang.String;@7b8a780a,[Ljava.lang.String;@7b8a780a,[Long.apache.spark.sql.sources.Filter;@73415c2d),DirectBigQueryRelation[fe-dev-sandbox.dkushari_ds.contents]),`dkushari_bq_test`.`dkushari_ds`.`contents`,Statistics(sizeInBytes=8.0 EiB, ColumnStat: N/A)
```

(8) PhotonRowToColumnar

Input [4]: [id#228, size#229L, binary#231, copies#232L]

(9) PhotonFilter

Input [4]: [id#228, size#229L, binary#231, copies#232L]

Arguments: isnotnull(id#228)

(10) PhotonShuffleExchangeSink

Input [4]: [id#228, size#229L, binary#231, copies#232L]

Arguments: hashpartitioning(id#228, 200)

(11) PhotonShuffleMapStage

Input [4]: [id#228, size#229L, binary#231, copies#232L]

Arguments: ENSURE\_REQUIREMENTS, [id=#636]

(12) PhotonShuffleExchangeSource

Input [4]: [id#228, size#229L, binary#231, copies#232L]

(13) PhotonShuffledHashJoin

Left keys [1]: [id#221]

Right keys [1]: [id#228]

Join type: Inner

Join condition: None

(14) PhotonProject

Input [7]: [repo\_name#217, path#219, id#221, id#228, size#229L, binary#231, copies#232L]

Arguments: [repo\_name#217, path#219, size#229L, copies#232L, binary#231]

(15) PhotonShuffleExchangeSink

Input [5]: [repo\_name#217, path#219, size#229L, copies#232L, binary#231]

Arguments: hashpartitioning(repo\_name#217, 200)

(16) PhotonShuffleMapStage

Input [5]: [repo\_name#217, path#219, size#229L, copies#232L, binary#231]  
Arguments: ENSURE\_REQUIREMENTS, [id=#644]

(17) PhotonShuffleExchangeSource

Input [5]: [repo\_name#217, path#219, size#229L, copies#232L, binary#231]

(18) PhotonSort

Input [5]: [repo\_name#217, path#219, size#229L, copies#232L, binary#231]  
Arguments: [repo\_name#217 ASC NULLS FIRST, size#229L DESC NULLS LAST]

(19) PhotonWindow

Input [5]: [repo\_name#217, path#219, size#229L, copies#232L, binary#231]  
Arguments: [repo\_name#217, path#219, size#229L, rank(size#229L)  
windowpecdefinition(repo\_name#217, size#229L DESC NULLS LAST,  
specifiedwindowframe(RowFrame, unboundedpreceding\$(), currentrow\$())) AS  
size\_rank#193, copies#232L, binary#231]

(20) PhotonShuffleExchangeSink

Input [6]: [repo\_name#217, path#219, size#229L, size\_rank#193, copies#232L,  
binary#231]  
Arguments: rangepartitioning(repo\_name#217 ASC NULLS FIRST, size\_rank#193  
ASC NULLS FIRST, 200)

(21) PhotonShuffleMapStage

Input [6]: [repo\_name#217, path#219, size#229L, size\_rank#193, copies#232L,  
binary#231]  
Arguments: ENSURE\_REQUIREMENTS, [id=#652]

(22) PhotonShuffleExchangeSource

Input [6]: [repo\_name#217, path#219, size#229L, size\_rank#193, copies#232L,  
binary#231]

(23) PhotonSort

Input [6]: [repo\_name#217, path#219, size#229L, size\_rank#193, copies#232L,  
binary#231]  
Arguments: [repo\_name#217 ASC NULLS FIRST, size\_rank#193 ASC NULLS FIRST]

(24) PhotonResultStage

Input [6]: [repo\_name#217, path#219, size#229L, size\_rank#193, copies#232L,  
binary#231]

(25) ColumnarToRow

Input [6]: [repo\_name#217, path#219, size#229L, size\_rank#193, copies#232L,

```
binary#231]
```

```
(26) AdaptiveSparkPlan
```

```
Output [6]: [repo_name#217, path#219, size#229L, size_rank#193,  
copies#232L, binary#231]
```

```
Arguments: isFinalPlan=false
```

```
== Photon Explanation ==
```

```
The query is fully supported by Photon.
```

## Observation:

- *Took around 11 minutes to run*
- *Pushdown to BQ did not happen*

## Query With Materialization:

### Query:

```
WITH joined_files AS (  
  SELECT  
    f.repo_name,  
    f.ref,  
    f.path,  
    f.mode,  
    f.id,  
    f.symlink_target,  
    c.size,  
    c.content,  
    c.binary,  
    c.copies  
  FROM  
    files WITH ('materializationEnabled' 'true') AS f  
    INNER JOIN contents WITH ('materializationEnabled' 'true') AS c  
      ON f.id = c.id  
)  
SELECT
```

```

repo_name,
path,
size,
RANK() OVER (PARTITION BY repo_name ORDER BY size DESC) AS size_rank,
copies,
binary
FROM
  joined_files
ORDER BY
  repo_name,
  size_rank;

```

## Explain Plan:

```

== Physical Plan ==
AdaptiveSparkPlan (7)
+- == Initial Plan ==
    Sort (6)
      +- Exchange (5)
        +- RunningWindowFunction (4)
          +- Sort (3)
            +- Exchange (2)
              +- EdgeBigQueryPushdownPlan (1)

(1) EdgeBigQueryPushdownPlan
Arguments: [SUBQUERY_46_COL_0#273, SUBQUERY_46_COL_1#275,
SUBQUERY_46_COL_2#285L, SUBQUERY_46_COL_3#288L, SUBQUERY_46_COL_4#287],
PreScala213BigQueryRDD[7] at RDD at PreScala213BigQueryRDD.java:40
External engine query: SELECT ( SUBQUERY_45.SUBQUERY_45_COL_0 ) AS
SUBQUERY_46_COL_0 , ( SUBQUERY_45.SUBQUERY_45_COL_1 ) AS SUBQUERY_46_COL_1
, ( SUBQUERY_45.SUBQUERY_45_COL_4 ) AS SUBQUERY_46_COL_2 , (
SUBQUERY_45.SUBQUERY_45_COL_6 ) AS SUBQUERY_46_COL_3 , (
SUBQUERY_45.SUBQUERY_45_COL_5 ) AS SUBQUERY_46_COL_4 FROM ( SELECT (
SUBQUERY_42.REPO_NAME ) AS SUBQUERY_45_COL_0 , ( SUBQUERY_42.PATH ) AS
SUBQUERY_45_COL_1 , ( SUBQUERY_42.ID ) AS SUBQUERY_45_COL_2 , (
SUBQUERY_44.ID ) AS SUBQUERY_45_COL_3 , ( SUBQUERY_44.SIZE ) AS
SUBQUERY_45_COL_4 , ( SUBQUERY_44.BINARY ) AS SUBQUERY_45_COL_5 , (
SUBQUERY_44.COPIES ) AS SUBQUERY_45_COL_6 FROM ( SELECT * FROM ( SELECT *

```



```
FROM (SELECT `repo_name`, `path`, `id` FROM
`fe-dev-sandbox`.`dkushari_ds`.`files` WHERE (`id` IS NOT NULL) ) AS
BQ_CONNECTOR_QUERY_ALIAS ) AS SUBQUERY_41 WHERE ( SUBQUERY_41.ID IS NOT
NULL ) ) AS SUBQUERY_42 INNER JOIN ( SELECT * FROM ( SELECT * FROM (SELECT
`id`, `size`, `binary`, `copies` FROM
`fe-dev-sandbox`.`dkushari_ds`.`contents` WHERE (`id` IS NOT NULL) ) AS
BQ_CONNECTOR_QUERY_ALIAS ) AS SUBQUERY_43 WHERE ( SUBQUERY_43.ID IS NOT
NULL ) ) AS SUBQUERY_44 ON ( SUBQUERY_42.ID = SUBQUERY_44.ID ) ) AS
SUBQUERY_45
```

## (2) Exchange

```
Input [5]: [SUBQUERY_46_COL_0#273, SUBQUERY_46_COL_1#275,
SUBQUERY_46_COL_2#285L, SUBQUERY_46_COL_3#288L, SUBQUERY_46_COL_4#287]
Arguments: hashpartitioning(repo_name#273, 200), ENSURE_REQUIREMENTS,
[plan_id=113]
```

## (3) Sort

```
Input [5]: [SUBQUERY_46_COL_0#273, SUBQUERY_46_COL_1#275,
SUBQUERY_46_COL_2#285L, SUBQUERY_46_COL_3#288L, SUBQUERY_46_COL_4#287]
Arguments: [repo_name#273 ASC NULLS FIRST, size#285L DESC NULLS LAST],
false, 0
```

## (4) RunningWindowFunction

```
Input [5]: [SUBQUERY_46_COL_0#273, SUBQUERY_46_COL_1#275,
SUBQUERY_46_COL_2#285L, SUBQUERY_46_COL_3#288L, SUBQUERY_46_COL_4#287]
Arguments: [repo_name#273, path#275, size#285L, rank(size#285L)
windowpecdefinition(repo_name#273, size#285L DESC NULLS LAST,
specifiedwindowframe(RowFrame, unboundedpreceding$(), currentrow$())) AS
size_rank#249, copies#288L, binary#287], [repo_name#273], [size#285L DESC
NULLS LAST], false
```

## (5) Exchange

```
Input [6]: [repo_name#273, path#275, size#285L, size_rank#249, copies#288L,
binary#287]
Arguments: rangepartitioning(repo_name#273 ASC NULLS FIRST, size_rank#249
ASC NULLS FIRST, 200), ENSURE_REQUIREMENTS, [plan_id=117]
```

## (6) Sort

```
Input [6]: [repo_name#273, path#275, size#285L, size_rank#249, copies#288L,
binary#287]
Arguments: [repo_name#273 ASC NULLS FIRST, size_rank#249 ASC NULLS FIRST],
true, 0
```

```

(7) AdaptiveSparkPlan
Output [6]: [repo_name#273, path#275, size#285L, size_rank#249,
copies#288L, binary#287]
Arguments: isFinalPlan=false

== Photon Explanation ==
Photon does not fully support the query because:
    Unsupported node: EdgeBigQueryPushdownPlan
    [SUBQUERY_46_COL_0#273, SUBQUERY_46_COL_1#275, SUBQUERY_46_COL_2#285L,
SUBQUERY_46_COL_3#288L, SUBQUERY_46_COL_4#287], PreScala213BigQueryRDD[7]
at RDD at PreScala213BigQueryRDD.java:40.

Reference node:
    EdgeBigQueryPushdownPlan [SUBQUERY_46_COL_0#273,
SUBQUERY_46_COL_1#275, SUBQUERY_46_COL_2#285L, SUBQUERY_46_COL_3#288L,
SUBQUERY_46_COL_4#287], PreScala213BigQueryRDD[7] at RDD at
PreScala213BigQueryRDD.java:40

```

## Observation:

- *Took around 23 minutes to run*
- *Pushdown to BQ did happen*

## Appendix:

### Things to try -

- SLA breach - is the workflow finishing within SLA?
- Try with materialization
- Check with different cluster configurations
  - Cluster DBR version (16.4 LTS+)
  - Cluster sizing